

Who Flips?

Self- and Cross-Model Counterarguments Reveal Answer Instability in LLMs

Nafiseh Nikeghbal^{1,2,3} · Amir Hossein Kargaran^{2,4} · Shaghayegh Koll^{1,2,3} · Jana Diesner^{1,2,3}

¹Technical University of Munich ²Munich Center for Machine Learning ³Munich Data Science Institute

⁴Ludwig Maximilian University of Munich · nafiseh.nikeghbal@tum.de



Paper



Code



Data

Evaluate stability
alongside accuracy.

01 / THE EXPERIMENT

Correct answers can flip under challenge.

01 FIRST ANSWER

Correct

02 COUNTERARGUMENT

A plausible case for
a wrong option.

Refusal → no challenge

03 ANSWER AGAIN

HOLD
correct

FLIP
any wrong

AFR = $\frac{\text{eligible challenges ending wrong}}{\text{all eligible challenges}}$


Answer flip rate is conditional: initially correct, with an available argument. Any final wrong option counts.

Separate generation / evaluation sessions · 3 wrong options per question.

02 / WHO FLIPS?

Answer stability varies across models.


Mean blind flip rates span
17.5–97.3% across the seven models.



03 / ARGUMENT LENGTH

Longer arguments are not always stronger.

Two Qwen models rise with longer arguments; GPT-5.1's decline is not statistically significant.



Four selected models. Length also changes the argument's content.

04 / SUBJECT DOMAINS


Domain matters.

Subject-level AFR spans
20.9–80.8%.

Selected extremes.
Means across models, lengths
and attribution conditions.

THREE LOWEST · STEM

Elementary math.
HS mathematics
College math.



THREE HIGHEST · HUMANITIES / SOCIAL SCI.

Moral disputes
Security studies
Professional law


20.9%
25.8%
30.4%

80.8%
80.6%
74.7%

05 / SOURCE → TARGET

Arguing ≠ resisting.

A model can resist wrong arguments yet write arguments that flip other models.




Outlined cells: same-model source.

06 / TWO DIFFERENT ROLES

Persuasive, yet hard to flip.

Epistemic porosity (EP): susceptibility as a target.

Epistemic authority (EA): efficacy as a source.




Off-diagonal column (EP) and row (EA) means.
Diagonal: EA = EP.

07 / STRONGER CHALLENGES

MaxFlip amplifies flips.

Selecting across source models raises AFR
by up to **23.6 percentage points**.




Versus same-model blind challenges.
No clear increase for GPT-5.1.

08 / HOW MAXFLIP WORKS

Pool. Test. Select.

Pool wrong-answer arguments across sources;
keep the most effective one for each question.



Schematic. Selection and evaluation use the same
model set; transfer to held-out models is untested.